pytorch/pytorch

Legend:

- Open
- Closed
- Closed (not planned)

1. `torch.jit.script` models with `Dict[str, Tensor]` return cannot be exported via `torch.onnx.export` without `dynamo=True`, and error message is unclear · Issue #155091 oncall: jit module: onnx triaged
2. Inconsistent behavior in `torch.export.export()` for empty forward functions with `*args` / `**kwargs`: one fails with `list index out of range`, the other succeeds · Issue #155190 oncall: pt2 oncall: export
3. `torch.export.export()` fails on GPU with LSTM model: "Cannot access data pointer of Tensor" · Issue #155309 good first issue oncall: pt2 export-triaged oncall: export
4. `torch.compile(..., mode="max-autotune")` fails with `TypeError: Expected a number but got Identity` under `torch.no_grad()` · Issue #155688 triaged oncall: pt2 module: dynamic shapes module: inductor
5. `nn.RNN(...).to('cuda')` fails with `cuDNN error: CUDNN_STATUS_BAD_PARAM` on GPU, but works on CPU · Issue #155798 module: cudnn module: cuda triaged
6. `torch.prod` or `torch.special.entr` triggers `CUDA driver error: invalid argument` on GPU unless kernel cache is cleared · Issue #156010 needs reproduction module: cuda triaged
7. `torch.export` fails with `KeyError` when `BatchNorm.running_mean` is read and modified, even when shape/value is unchanged · Issue #156167 triaged oncall: pt2 export-triaged oncall: export
8. `torch.compile(..., mode="max-autotune", dynamic=True)` causes small but nonzero output mismatch with `nn.Conv2d` compared to eager output · Issue #156301 triaged oncall: pt2 topic: fuzzer
9. `torch.compile(fullgraph=True, options=...)` fails with `NoValidChoicesError` on simple `Conv2d` model, but gives no actionable trace · Issue #156304 triaged oncall: pt2
10. `torch.jit.load()` crashes with C++ `c10::Error` if file system is full (no graceful handling) · Issue #156955 oncall: jit module: error checking actionable module: edge cases
11. Inconsistent behavior between eager and compiled mode for `F.conv_transpose2d` · Issue #157909 oncall: pt2
12. Question: Is it expected that `QuantStub` and `DeQuantStub` are skipped in `torch.compile`? · Issue #157998 oncall: pt2
13. `torch.compile` cannot compile a model with a basic `LSTM`, even on latest main · Issue #158007 module: rnn triaged oncall: pt2 module: dynamo
14. `torch.compile` leads to numeric mismatch compared to `float64` eager baseline · Issue #158530

triaged oncall: pt2 module: inductor

15. Bug or Precision Tradeoff? Small `torch.compile` Mismatch with `max-autotune-no-cudagraphs` on CUDA · Issue #158544 triaged oncall: pt2 module: inductor
16. Inconsistent Error Message for Cross-Device Input in torch.compile · Issue #159133 module: error checking triaged oncall: pt2 module: dynamo
17. `torch.compile(nn.ModuleList())` changes `bool()` behavior · Issue #159139 triaged oncall: pt2 module: dynamo
18. Bug: `torch.compile` triggers C++ compile error due to conflicting declaration in generated `.cpp` code · Issue #159245 needs reproduction oncall: pt2 oncall: cpu inductor
19. `torch.compile` crashes when accessing a `weakref.proxy` inside a compiled method · Issue #159258 triaged oncall: pt2 module: dynamo
20. torch.cond triggers shape mismatch with linear layer when nested in custom model · Issue #159353 triaged oncall: pt2 module: higher order operators
21. `torch.compile` fails with `InternalTorchDynamoError` when module defines attribute named `parameters = None` · Issue #159460 triaged oncall: pt2 module: dynamo
22. `torch.cond()` behaves inconsistently when using symbolic predicate · Issue #159852 triaged oncall: pt2 module: higher order operators module: pt2-dispatcher
23. `torch.compile` crashes on model exported via `torch.export`, with lifted constant from cache tensor · Issue #160066 oncall: pt2 oncall: export
24. `torch.export.export` fails when model is in `eval()` mode due to `torch.typename` being skipped by Dynamo · Issue #160241 triaged oncall: pt2 module: fakeTensor module: dynamo oncall: export
25. `torch.compile` produces output mismatch vs eager float64 for some seeds · Issue #160242 high priority triaged oncall: pt2 module: inductor
26. `torch.compile` backward pass fails with `AssertionError` in Inductor C++ codegen when model returns tuple outputs · Issue #160391 oncall: pt2 oncall: cpu inductor
27. `torch.compile` causes segmentation fault with simple `ModuleList` model after multiple compile calls · Issue #160399 oncall: pt2 oncall: cpu inductor
28. `torch.compile` (Inductor) segfaults on `torch.slice_copy` with huge step; eager and `aot_eager` OK · Issue #160868 high priority triage review triaged actionable oncall: pt2 module: inductor oncall: cpu inductor
29. `torch.compile` (Inductor) fails on MKL-DNN tensor densify (`to_dense`) with "Cannot access data pointer of Tensor that doesn't have storage" · Issue #160873 oncall: pt2 oncall: cpu inductor
30. `torch.export` with invalid dynamic_shapes expressions raises `AttributeError: 'One' object has no attribute 'name'` instead of a clear error · Issue #161902 triaged oncall: pt2 oncall: export
31. `torch.compile` fails with `aten.complex.default` size/stride assertion when using STFT output recombined into complex tensor · Issue #161906 triaged oncall: pt2 module: inductor
32. `torch.compile` produces inconsistent results vs eager (FP64 baseline) with data-dependent indexing and reciprocal ops · Issue #161933 triaged oncall: pt2
33. `torch.compile` fails on `pad_packed_sequence` after GRU with `PackedSequence` (FakeTensor "data is not allocated" error) · Issue #162374 triaged oncall: pt2 module: dynamo
34. `torch.compile` fails with `pack_padded_sequence` + `pad_packed_sequence` (FakeTensor error) · Issue

#162375 triaged oncall: pt2 module: dynamo

35. [ONNX] Regression in 2.8.0: exporting LSTM with `nn.utils.weight_norm` aborts in legacy TorchScript exporter (works in 2.7.1) · Issue #162376 module: onnx triaged
36. [ONNX][Legacy Exporter] Segmentation fault when `dynamic_axes` keys don't match input/output names (should error, not crash) · Issue #163033 module: onnx triaged OSS contribution wanted
37. [compile/inductor][functionalization] `prims.broadcast_in_dim` triggers alias-annotation error under `torch.compile` (eager OK) · Issue #163037 triaged oncall: pt2 module: aotdispatch module: pt2-dispatcher
38. [torch.export.load] Passing an `ExportedProgram` raises `AttributeError: 'ExportedProgram' object has no attribute 'tell'` instead of a clear `TypeError` · Issue #163040 triaged oncall: pt2 oncall: export
39. [compile/aot-autograd] `torch.autograd.grad` w.r.t. an intermediate tensor works in eager but fails under `torch.compile` ("tensor not used in the graph") · Issue #163048 triaged oncall: pt2 module: aotdispatch module: pt2-dispatcher
40. [torch.export] `AssertionError: assert isinstance(a, FakeTensor)` in `_free_unbacked_symbols_with_path` when exporting model that uses `torch.func.jvp` + `functional_call` + dict params · Issue #163051 oncall: pt2 oncall: export
41. torch.compile: Segfault or "Unhandled FakeTensor Device Propagation" when compiling a Module on CPU, moving it to CUDA, recompiling, then calling with a CPU input · Issue #163064 triaged oncall: pt2 module: fakeTensor module: dynamo module: guards module: pt2-dispatcher
42. [torch.compile] TransformerEncoder with boolean src_key_padding_mask fails under Dynamo fullgraph mode · Issue #163640 module: nn module: error checking triaged oncall: pt2
43. [inductor][pattern_matcher] RuntimeError in joint graph passes when compiling simple `randperm` + advanced indexing pipeline · Issue #166604 triaged oncall: pt2 module: inductor inductor_pattern_match
44. `torch.export.export(..., dynamic_shapes=..., strict=False)` crashes with `AssertionError` in `symbolic_shapes.prettify_results` (`assert op == "=="`), triggered by `torch.stft` + dynamic time dim · Issue #171673 module: bootcamp oncall: pt2 oncall: export
45. `torch.export.export(torch.compile(model), ...)` fails with Inductor `fake mode mismatch` assertion (`detect_fake_mode`) · Issue #171674 oncall: pt2 oncall: export
46. `torch.compile` with Inductor (dynamic=True, fullgraph=True) fails on ViT-like block with SymPy Relational `__bool__` error in `ir.get_stride_order` · Issue #172204 triaged oncall: pt2 module: dynamic shapes module: inductor
47. Eager buffer mutation after `torch.compile` + backward + `torch.export.export` hits `_map_assigned_buffer_to_proxy` `FakeTensor` assertion · Issue #172213 triaged oncall: pt2 module: aotdispatch module: pt2-dispatcher
48. Deprecated `torch._export.aot_load()` crashes with segmentation fault when loading compiled AOT artifacts with constants · Issue #172739 high priority module: crash oncall: pt2 oncall: export module: aotinductor
49. `torch.linalg.ldl_solve` triggers `double free or corruption` in eager mode (CPU) · Issue #172744 high priority triage review module: crash triaged module: linear algebra
50. `torch.compile` fails in FakeTensor meta path: `Cannot view ... strides ... as (1, 2048)` while

eager works · Issue #172830 triaged oncall: pt2 module: fakeTensor module: dynamic shapes module: decompositions module: pt2-dispatcher

tensorflow/tensorflow

1. Inconsistent output after model reload when using @tf.function on `call` · Issue #98230 stat:awaiting response type:bug stale comp:keras TF 2.19
2. Inconsistent model output after save/load when using `@tf.function(jit_compile=True)` · Issue #98233 stat:awaiting response type:bug stale comp:ops comp:xla TF 2.19
3. Invalid `class_weight` silently ignored in `model.fit()`, no error raised · Issue #98283 type:bug comp:keras comp:ops awaiting PR merge TF 2.19
4. `tf.function(jit_compile=True)` fails with GRU/LSTM (CuDNNRNNV3) on XLA GPU JIT: "No registered 'CudnnRNNV3' OpKernel for XLA_GPU_JIT" · Issue #99385 stat:awaiting response type:bug stale comp:keras comp:xla TF 2.19
5. tf.image.combined_non_max_suppression aborts process on wrong-rank inputs (should raise Python exception instead) · Issue #99583 type:bug comp:ops awaiting PR merge TF 2.19
6. [GPU] CUDA_ERROR_ILLEGAL_ADDRESS / “No algorithm worked” in Conv2D after full spatial self-attention at 224×224 (N≈50k → N²≈2.5B). Looks like index/size limit; should fail earlier with a clearer error · Issue #99631 type:bug comp:gpu TF 2.19
7. [GPU][BiasAdd] Kernel launch count overflow with large Dense/bias_add (negative `work_element_count` → abort) · Issue #99894 type:bug comp:gpu TF 2.19
8. XLA GPU JIT: keras.layers.RandomRotation fails with "No registered 'ImageProjectiveTransformV3' OpKernel for XLA_GPU_JIT" (TF 2.20.0, Keras 3.11.3) · Issue #100865 type:bug comp:keras comp:xla TF 2.19
9. XLA GPU JIT: Keras LSTM hits "No registered 'CudnnRNNV3' OpKernel for XLA_GPU_JIT" (TF 2.20.0, Keras 3.11.3) · Issue #100872 stat:awaiting response type:bug stale comp:keras comp:xla TF 2.19
10. XLA GPU JIT: tf.random.poisson lowers to RandomPoissonV2 which has no XLA_GPU_JIT kernel (fails under tf.function(jit_compile=True)) · Issue #100874 type:bug comp:gpu comp:xla TF 2.19
11. [TF/Keras][Conv3D+MaxPool3D] Zero-sized spatial dims cause fatal cuDNN abort (should raise Python InvalidArgumentError instead) · Issue #101409 type:bug comp:keras TF 2.19
12. [GPU][UnsortedSegmentProd] GPU kernel assertion on large num_segments (overflow in launch config) · Issue #102383 type:bug comp:ops comp:gpu TF 2.19
13. `tf.function(..., jit_compile=True)` inside `tf.GradientTape` fails with `UNIMPLEMENTED: TensorList crossing XLA/TF boundary`, followed by BFCAllocator fatal abort · Issue #108080 type:bug type:support comp:tf.function

apache/tvm

1. [Bug] Segfault when instantiating abstract `SearchStrategy()` in `TuneContext` · Issue #18268 type: bug needs-triage
2. [Bug] [Relax][DLPack] 0-D `bool` from PyTorch via DLPack mis-typed at VM entry → dtype check failure · Issue #18315 type: bug needs-triage

3. [Bug] [Crash][FFI/LLVM] Segfault when importing TVM after PyTorch/Transformers — LLVM static init in `COFFDirectiveParser` (`OptTable::buildPrefixChars`) · Issue #18316 type: bug needs-triage
4. [Bug][Relax][Torch] Segfault in from_exported_program when model returns (Tensor, None) tuple · Issue #18337 type: bug needs-triage
5. [Bug][Relax][Torch] from_exported_program fails on unbind producing single-element tuple · Issue #18338 type: bug needs-triage
6. [Feature Request][Relax][Torch] Add support for torch.mm (mm.default) in from_exported_program · Issue #18339 type: bug needs-triage
7. [Feature Request][Relax][Torch] Add support for nn.LSTM (lstm.input) in from_exported_program · Issue #18340 type: bug needs-triage
8. [Bug][relax.frontend.torch] Segfault in relax::Tuple when importing/building a torch.export program that uses batch-wise advanced indexing write (aten.index_put_) · Issue #18355 type: bug needs-triage
9. [Bug] [feature][relax.frontend.torch] from_exported_program rejects fused GRU (aten::gru.input) with "Unsupported function types ['gru.input']" · Issue #18356 type: bug needs-triage
10. [bug][relax.frontend.torch] from_exported_program KeyError for non-persistent buffer bert.embeddings.position_ids (HuggingFace BERT via torch.export) · Issue #18357 type: bug needs-triage
11. [bug/feature][relax.frontend.torch] from_exported_program rejects randn.default (blocks repro that stresses advanced-indexing write + tuple output) · Issue #18362 type: bug needs-triage
12. [bug][relax.frontend.torch] FFI segfault in tvm::relax::Tuple::Tuple when importing torch.export graph with 4D advanced-indexing write (aten.index_put_) and tuple outputs · Issue #18363 type: bug needs-triage
13. [feature][relax.frontend.torch] Missing coverage for STFT+RNN pipeline: 'rnn_tanh.input', 'real.default', 'imag.default', 'unfold.default', 'fft_fft.default' in from_exported_program · Issue #18364 type: bug needs-triage
14. [feature][relax.frontend.torch] Missing lowering for anti-aliased bilinear resize (aten::_upsample_bilinear2d_aa.vec) in from_exported_program · Issue #18365 type: bug needs-triage
15. [Bug] [relax][torch] from_exported_program segfault with exported MHA using eq(0)/expand mask + in-place masked_fill_ (get_attr lifting warning from PyTorch) · Issue #18407 type: bug needs-triage
16. [Bug] `relax.frontend.torch.from_exported_program` aborts on sparse CSR buffer (`layout_impl is only implemented for TensorImpl subclasses`) · Issue #18648 type: bug needs-triage
17. [Bug] Segfault in LLVM COFFOptTable when combining TVM CUDA target with PyTorch Lightning + torchmetrics in the same process · Issue #18651 type: bug needs-triage
18. [Bug] Segmentation fault when using TVM together with `transformers` (flan-t5-base + bfloat16 + `use_cache=True`) · Issue #18653 type: bug needs-triage
19. [Bug] Segmentation fault when importing `torchmetrics` after creating a TVM CUDA target (LLVM initialization / COFF OptTable) · Issue #18655 type: bug needs-triage

Oneflow-Inc/oneflow

1. [bug] `F.adaptive_max_pool2d` with `non-4D` input crashes the process (C++ CHECK abort) instead of raising a Python error · Issue #10669 bug community

2. [bug] MaxUnpool2d with invalid indices (-1) crashes process instead of raising Python error · Issue #10670 bug community
3. [bug] CrossEntropyLoss with wrong target shape aborts process (C++ CHECK) instead of raising Python error · Issue #10671 bug community
4. [bug] flow.inverse on singular matrix triggers FATAL abort instead of raising RuntimeError · Issue #10672 bug community
5. [bug][nn.Graph] Accessing submodules/buffers/params via user dict leads to segfault (should hard-error with clear message) · Issue #10673 bug community
6. [bug][nn.Graph] deepcopy(Graph) raises KeyError('_backward_hooks') (should be unsupported with clear TypeError or deep-copiable) · Issue #10674 bug community
7. [bug][nn.Graph] flow.full with 0-D Graph proxy tensor as fill_value crashes with "This is a oneflow bug" · Issue #10675 bug community
8. [bug][nn.Graph] stft with registered window and pad_mode='reflect' crashes in Graph mode ("This is a oneflow bug") · Issue #10676 bug community
9. [bug][nn.Graph] pack_padded_sequence → LSTM → pad_packed_sequence crashes during Graph build ("This is a oneflow bug") · Issue #10677 bug community
10. [bug][nn.Graph] In-place advanced indexing assignment (y[idx0, idx1] = v) crashes during Graph build ("This is a oneflow bug") · Issue #10678 bug community
11. [bug][graph] Advanced indexing assignment with negative indices crashes in Graph mode (RuntimeError: This is a oneflow bug ...) · Issue #10679 bug community
12. [bug][graph] Segfault when reading a registered buffer via module._buffers["0"] inside forward (eager OK) · Issue #10680 bug community
13. [bug][nn.CrossEntropyLoss] Wrong-shaped target causes internal CHECK failure (is_initialized()) and abort instead of shape-mismatch error · Issue #10681 bug community